

4 The Relationship between Surprisal and RTs (60 points)

Problem Setup

Smith and Levy (2013) explored the shape of the relationship between a word's predictability in context and how long, on average, a comprehender spends on reading that word. In particular they were interested in the hypothesis that a word's average reading time might be linear in the SURPRISAL of the word (Hale, 2001; Levy, 2008):

$$\text{surprisal}(w_i|\text{Context}) = \log \frac{1}{P(w_i|\text{Context})}$$

or in some other function of the word's conditional probability. In this assignment, you will investigate this question yourself, using one of the datasets analyzed by Smith and Levy.

The Smith and Levy (2013) self-paced reading dataset derives from each subject in the experiment reading a number of several-hundred-word passages selected from the Brown corpus (Kučera & Francis, 1967). This dataset is presented in tabular format, with one reading-time measurement per row and with the following columns:

- The **word** that was read;
- A **code** that uniquely identifies the word/context pair;
- The identifier for the **subject** in the experiment from which the reading-time measurement was taken;
- **text_id**, which is the identifier for the text from the Brown corpus that was being read;
- **text_pos**, which is the word number in the text selection that was being read;
- **word_in_exp** is the word number in the experiment for the particular subject;
- The **time** in milliseconds that the word in question was visible on-screen during the subject's reading. (Remember, in self-paced reading this is the time elapsed between the subject pressing a button/key to reveal the word, and the subject pressing the button/key again to mask that word and reveal the next word.)

Your task

Your task for this assignment is to reproduce the main result of Smith and Levy (2013), using the output of a 5-gram model, trained on part of the Penn Treebank, a commonly used dataset in natural language processing research.

The notebook for this assignment can be found [here](#). Please carefully read and follow all instructions in the notebook. You will need to write code or textual responses in every place marked as **TODO**.

Obtaining n -gram surprisals/RTs

We will provide the reading time (RT) data in a file called `brown_rts.csv`. We will also provide per-token surprisal data from a pre-trained n -gram model (a Kneser-Ney-smoothed 5-gram model) in a file called `ngram_surprisals.tsv`. Your first step is to combine these two files into a format where you can analyze the relationship between reading time and word surprisal. Please see the Colab notebook for detailed instructions.

Before you can do this, in the process of estimating word probabilities, you will need to attend carefully to issues of **word tokenization**. The way words are presented to humans in the reading experiment does not match the way words are separated and processed by the n -gram model. In the self-paced reading experiment, every button press presented a single contiguous string of non-space ASCII characters, so the dataset reflects tokenization by spaces. In the n -gram training data, by contrast, punctuation characters are tokenized out from the rest of word strings.

You will also need to handle out-of-vocabulary (OOV) items. Since not all words in the self-paced reading dataset appear in the n -gram's vocabulary, some of the surprisal values will correspond to unknown (“unk”) tokens. In your writeup, describe what preprocessing steps you performed.

Statistical Analyses

Once you have aligned the surprisals and the RTs, your job is to visualize and statistically analyze their relationship. You will examine and compare the relationship of RTs both to surprisals and to raw probabilities.

Specifically, for each metric in [surprisal, raw probability]:

- Fit a linear regression model to predict RTs from the metric. You should report the coefficient for the metric term (slope) and a corresponding t -score and p -value (to determine whether it is significantly different from 0), as well as an R^2 -score (the coefficient of determination) of the model.
- Draw metric-RT scatterplot with best-fit line, **without** binning RT values; and
- Draw metric-RT scatterplot with best-fit line, **with** binning RT values.

Interpret the results Does the univariate analysis support the hypothesis of a linear relationship between word surprisal and word reading time? Is that hypothesis better or worse than an alternative hypothesis of a linear relationship between *raw word probability* and word reading time? Are there other alternative hypotheses that might be even more compelling given the data?

Multiple regression analysis: The aim of this analysis is to assess the surprisal-RTs association while controlling *word-length* and *word log-frequency*. First, you will be asked to extract these two variables by yourself (see detailed instructions in the Colab notebook). Then, you should fit a multiple linear regression model to predict RTs from *surprisal + word length + word log frequency*, interpret the model's results, and compare them to the results from analysis 1. The new model should take the form:

$$\text{RT} = \beta_0 + \beta_1 \text{surprisal} + \beta_2 \text{length} + \beta_3 \log \text{ frequency}.$$

Interpret the results How does the *surprisal* coefficient of this model compare to the surprisal coefficient in the univariate model? Does your conclusion regarding the effect of *surprisal* on RTs from the univariate analysis still hold?

What you need to turn in


Your full Colab notebook copy, including not only code, cell output, graphs, and quantitative analyses, but also a verbal description of what you did and your interpretation of the results. You can save your notebook as a PDF and submit this file as your write-up.

Notes and suggestions

- It is common in self-paced reading to exclude “outlier” reading times that might reflect non-linguistic processing effects, such as the participant taking a break mid-sentence. You can find the outlier-removal policy used in Smith and Levy (2013) in the *Materials and Methods* section of the paper. You can use this policy, an alternative policy, or simply conduct no outlier removal. Be clear about the policy you take and the motivation for it, and take it into account when you provide your interpretation of the results.
- For data visualization, you can use whatever software you are comfortable with. If you have not made data plots before, we recommend using `matplotlib`, a Python visualization library. If you don't have it installed already, you can install it via `pip` with the command `pip install matplotlib`. Here is an example of creating a scatterplot with lines between the points and axis labels:

```
import matplotlib.pyplot as plt
x = [1, 2, 3, 4] # x coordinates for the example plot
y = [2, 3, 4, 6] # y coordinates for the example plot
```

```
plt.scatter(x, y) # Create a scatterplot with given x and y coordinates
plt.plot(x, y) # Draw lines between the points
plt.xlabel("X axis")
plt.ylabel("Y axis")
plt.show() # Display it on your screen
plt.savefig("filename.pdf") # Save the result to filename.pdf
```



References

- Hale, J. (2001). A probabilistic Earley parser as a psycholinguistic model. *Proceedings of the Second Meeting of the North American Chapter of the Association for Computational Linguistics*, 159–166.
- Kučera, H., & Francis, W. N. (1967). *Computational analysis of present-day American English*. Providence, RI: Brown University Press.
- Levy, R. (2008). Expectation-based syntactic comprehension. *Cognition*, 106(3), 1126–1177.
- Smith, N. J., & Levy, R. (2013). The effect of word predictability on reading time is logarithmic. *Cognition*, 128(3), 302–319.